
Towards Reactive Information Systems and their Services

MASTER THESIS

Author:
Dominic BOSCH

Supervisors:
Prof. Dr. Helmar BURKHART
Dr. Martin GUGGISBERG

June 24, 2014



Abstract

The Web is a rapidly growing information universe, consisting of Information Systems that provide access over heterogeneous services. The majority of those Information Systems are dynamic and changes in their Information Space can be regarded as events, when aiming to detect them. Moreover, such changes are also actions, when imposed onto the Information Space. If appropriate services exist to access such an Information System for read and write operations, we are able to orchestrate it. By adopting the Event-Condition-Action (ECA) paradigm to Information Systems, we are able to introduce an event-based conceptual model, which allows the detection of events and the dispatching of actions according to predefined rules, thus imposing reactivity on top of or between Information Systems. Current approaches that use the ECA paradigm focus on action imposition on local storage, while we aim to impose actions on the heterogeneous services of existing Information Systems. This model is not limited to the Web, but can include any in a way accessible Information Systems over its services.

In our work we introduce a prototype system, which uses the Web's programmability to impose reactivity on top of it. We also underline the importance of the, currently little, support of Web services for event callback addresses, the so called Webhooks. They are the only way for real-time event delivery to remote systems and free them from expensive polling for changes. Through our prototype it is possible to orchestrate Web services based on an Event-Driven Architecture (EDA), a method which pushes towards the vision of real-time reactive Information Systems. We list some example use cases for our conceptual model, as well as use cases that have been implemented in our prototype system.

Acknowledgment

I would like to express my deep gratitude to my master thesis advisor, Prof. Dr. Helmar Burkhart, for his patience and constant encouragement to keep going further. I also want to thank Dr. Martin Guggisberg for his ever helpful advices, hints to existing technologies and guidance through the whole process of this thesis. Besides my advisors, I would like to thank the members of the research group: Danilo Guerrero, Antonio Maffia, Alexander Gröflin and Robert Frank for their help, many interesting discussions and inputs. My thanks also go to Manus Bonner and Benedikt Willi for reading through this and coming up with well formed sentences and technical corrections. Finally I would like to thank my family; my parents Silvia and Peter Bosch, my sister Svenja, her husband Andreas Buser, their seven months old son Nico and my girlfriend Kathrina for their inspiring ideas, the joyfulness, irresistible smiles and financial support they gave me on my journey.

Contents

Glossary	vii
Acronyms	ix
1 Introduction	1
2 Related Work	3
2.1 Data and Functionality Providers on the Web	3
2.2 Reactivity through Event-Condition-Action Rules	7
3 Conceptual Model for Reactive Information Systems and their Services	12
3.1 From Physical Events to Virtual Events	13
3.2 Capturing Events from Information Systems	14
3.3 Event Pattern Detection	14
3.4 Imposing Reactivity to Information Spaces	15
4 Use Cases for Reactive Information Systems	16
4.1 Reacting on changes in the World Wide Web	16
4.2 Enhance existing Web Applications	16
4.3 Service Functionality and Availability Checking	18
4.4 Exploiting the Web of Things	19
4.5 Multi-Sourced Bad Weather Prediction	19
5 Prototype System	21
5.1 Architecture	21
5.2 A Rule Language for the Prototype System	27
5.3 Prototype Use Case Implementations	28
5.4 Web Application Development	28
6 Conclusions & Future Work	32
Bibliography	33
Appendices	37

List of Figures

1.1	Users exploit the Web's Programmability through a Reactive Entity in the Web	2
2.1	Number of registered APIs in the ProgrammableWeb directory by date	4
3.1	Reactivity imposed on Information Systems and their Information Spaces over Services	12
3.2	Conceptual Model for Reactive Information Systems and their Services	13
3.3	Information Systems providing access to the Web of Things over their Services	15
4.1	Enrich CMS Post with Remote Knowledge Data	17
4.2	Create course resources at semester start	17
4.3	Create course resource for registered student	18
4.4	Notify student before exercise due date	18
4.5	Test proper Service Functionality and Availability	19
4.6	Measurements on Server Failure	20
5.1	Prototype System Architecture	22
5.2	UC Binder Annotation	29
5.3	Synchronous Function Call	29
5.4	Asynchronous Function Call	30
5.5	Closure Scope and referenced context	30
5.6	Closure context changes in a loop	31
5.7	Closure Context Shielding	31

List of Tables

2.1	Key Properties of existing Rule Languages	11
-----	---	----

Listings

5.1	Event Trigger code to poll Email Yak RESTful Web service for new Mails; written in CoffeeScript	25
5.2	Action Dispatcher code to store a new content on the ProBinder RESTful Web service; written in CoffeeScript	26
5.3	Rule Example expressed in JSON	27
5.4	Example Phrase in Prototype Rule Language	28
5.5	Extended Backus-Naur Form of Prototype Rule Language Syntax	28
5.6	JavaScript Closure Context Shielding	31

Glossary

Information Space *"[...] is a set of concepts and relations among them held by an Information System. Information Space is produced by a set of known procedures, and is changed through intentional manipulation of its content"[27].* i, iv, 2, 3, 12–16, 21, 24

Information System is a network of software and hardware components that support collection, filtering, storing, processing and distribution of data. i, iv, 1–5, 7–9, 12–16, 21, 24, 25, 32, 33

Mashup A Web Application that weaves two or more different Web APIs together to provide a new perspective on data. 4, 6

Semantic Web Tim Berners-Lee's vision of the machine-readable Web through standard data formats and semantic metadata descriptions on top of the resources via RDF which turn the Web into a structured data collection. 1, 5, 8

Web API An Application Programming Interface to either a Web service or the browser, used for application to application communication. 3, 4, 6–8, 23, 24, 32, 33

Web Application An application which runs in the browser. Often a user interface to an application sitting on a server. Single Page Applications (SPA), a subset of the Web Applications, access Web Resources over asynchronous calls to the server while the user is interacting with the application. 16, 17, 32

Web of Things An evolution of the Internet of Things, which describes the integration of smart things (e.g. sensors, embedded devices or digitally enhanced objects) into the Internet. The Web of Things is the adoption of the REST architectural style to the smart things in order to enable uniform access to these loosely coupled entities. iv, 1, 13, 15, 33

Web Resource Anything in the Web which can be identified, addressed and handled. Identification and addressing is often done over URIs. Web Resources and their semantic properties are described using RDF in the Semantic Web. 1–3, 5–8, 10–12, 14, 16, 17, 23, 25, 32, 33

Web Service A collection of SOAP related Web service (note the lower-case word service) standards which are widely adopted and developed in the industry. Also called "WS-*" Web Services or the Big Web Services. 3, 4, 23

Web service An interface for communication between applications over a network. They can provide access to and control over Web Resources. Web services are also called services on the Web or just services. i, 3–7, 13, 23, 24, 32, 33

Webhook A server-side Web API which accepts a URI that is used as a callback to push events to an external Web Resource. i, 7, 14, 21, 22, 24, 25, 32, 33

World Wide Web Tim Berners-Lee's vision of interlinked hypertext documents which are accessed over the Internet via browser and allow the navigation through a global information universe. iii, 3, 5, 14, 16, 24

Acronyms

- CED** Complex Event Detection. 10, 14
- CEP** Complex Event Processing. 8, 10, 14, 33
- CMS** Content Management System. iv, 16, 17
- CORBA** Common Object Request Broker Architecture. 5, 9
- DOM** Document Object Model. 8, 11
- ECA** Event-Condition-Action. i, iii, 3, 7, 8, 10, 14, 15, 21, 25, 27, 33
- EDA** Event-Driven Architecture. i, 7, 9, 10, 15, 21
- ESB** Enterprise Service Bus. 9, 10
- HTML** Hypertext Markup Language. 32
- HTTP** Hypertext Transfer Protocol. 5, 33
- IaaS** Infrastructure as a service. 3
- IDL** Interface Definition Language. 5
- IIOP** Internet Inter-ORB Protocol. 5
- JSON** JavaScript Object Notification. vi, 7–9, 11, 21–23, 26, 27
- KR** Knowledge Representation. 8
- KRE** Kinetic Rules Engine. 9, 11
- KRL** Kinetic Rule Language. 9–11
- ORB** Object Request Broker. 5
- PaaS** Platform as a service. 3
- RDF** Resource Description Framework. 8, 11, 33
- RDFTL** RDF Triggering Language. 8, 11
- REST** Representational State Transfer. 3, 5, 7, 23, 32, 33

REVERSE Reasoning on the Web with Rules and Semantics. 8

RPC Remote Procedure Call. 4, 5

RuleML Rule Markup Language. 8, 9, 11

SaaS Software as a service. 3

SOA Service-Oriented Architecture. 3, 4, 21

SOAP Simple Object Access Protocol. 3–5, 23

URI Uniform Resource Identifiers. 5, 7, 24, 25

WSDL Web Service Description Language. 4, 5

XML Extensible Markup Language. 4, 7–9, 21

XML-RPC XML - Remote Procedure Call. 4

Chapter 1

Introduction

Information Systems are omnipresent in today's world, and being connected to the Web is more a common attribute than a special quality among them. As a consequence, the Web is an ever growing institution in all aspects that it covers. The number of data and functionality providing Information Systems is growing in the whole spectrum from bigger computing centers down to smaller devices. Computing centers are growing in size and quantity and they allow for massive amounts of data to be stored and accessed, moreover they also enable the construction and offering of more complex functionality. At the same time, an increasing number of ever smaller devices also provides more Information Systems attached to the Web. Many of them are offering access to the Web itself, granting even more devices access and thus leverage the effect of the growing Web, e.g. mobile phones can act as a hotspot to grant Web access to other devices over WiFi. A recent observed trend is all the smart things, which have access to the Web and start to form the Web of Things. Today, these smart things can be everything from a temperature sensor to all the electronic devices within a house. They do not only provide sensor data but they can also be controlled over the Web. All these different types of services available on the Web make it a heterogeneous collection of Information Systems and their services. Great efforts are made to turn them into uniformly accessible Web Resources, e.g. the Semantic Web is a widely supported initiative towards a machine-readable, structured and semantically described Web.

Confronted with this rapid growth of the Web, an increasing number of human beings is exposed to it in their daily life, and they get literally flooded with informations and means to retrieve or process them. Even though users have access to so many Web Resources, they often lack the knowledge, necessary time or right approach to weave them together. Great value would be added for them if they could automatically get appropriate informations, in the right moment and in a condensed matter that supports them best. They should be able to automate tedious tasks, e.g. detecting relevant changes in their preferred Web Resources and react on behalf of such changes. This requires the identification of and filtering for user-relevant changes, appropriate timing, assembly and finally the placement of the outcome in the user's preferred Web Resources.

With the many existing Information Systems and their services on the Web, users don't want to be bound to specific ones for certain tasks, as it is often the case nowadays. They want to use the functionality or data of their preferred ones, which helps them best to fulfill their needs. Hence, users should have to be able to create their own specific but still flexible Web Resource orchestrations. Since the need of users to orchestrate different services has gotten a lot of attention, some of the Information Systems on the Web offer ways to spread their data to others, but in a limited way. For example it is common for social network applications to push user-specific notifications to other social networks, e.g. singing in at a place in Foursquare

can also be posted directly to the Facebook timeline. Because of existing limitations, such as customizability or action imposing on Information System of their choice, users still end up mixing data and functionality from different Web Resources by hand, which often means to execute similar tasks repeatedly by hand. Moreover the manual reaction on changes is deferred because the changes are not detected in real-time by the user or because the user is not able to react in a timely fashion.

Since data and functionality already exist in the Web, the users are theoretically enabled to automate their work to some extent by orchestrating those Web Resources. Even though the access to resources gets simpler, the average user is still not capable to fully exploit the Web's full potential. Another challenge is, that often a lot of effort has to be made, in understanding how the specific service works, before it can be fully exploited. There is a lot of research that goes towards an easy to orchestrate the Web, but they are either often complicated to wield themselves, mere data copy tasks or static Web Resource compositions. Our goal is to enable user-defined resource orchestrations, and still exploiting their full potential by not limiting the set of their functionality and thus going towards a reactive Web.

A big part of the data, that becomes available to the users, is short-lived data that corresponds to state changes, which are changes in the Information Spaces and can be modeled as events for detection or actions for imposition. In this thesis we introduce an event-driven conceptual model that uses the programmability of Information Systems and their services in order to impose reactivity between them. By regarding the whole Web as an Information Space, in which we listen for triggered events and on which we impose actions as a result of user-defined rules, we are able to model a real-time reactive Web, as shown in Figure 1.1. Such a personalized reactivity allows the orchestration of the Web and a tool to govern its data flood by automating tedious tasks. Current Web Resource orchestrations concentrate on data flow rather than on event flow, which are mere copy/paste tasks of data than smart reactivity. This makes us believe that our event-based conceptual model can overcome certain shortcomings of the existing approaches and provides a step towards the real-time reactive Web.

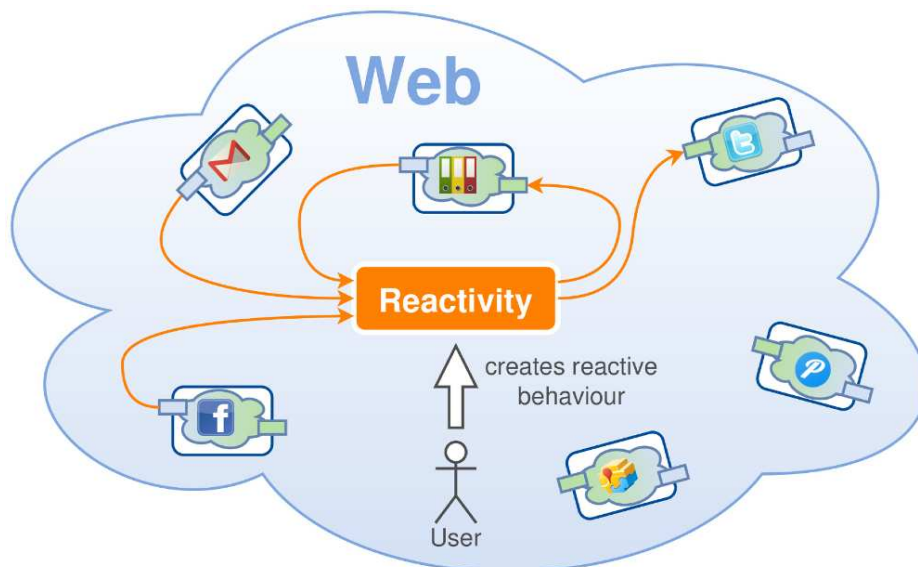


Figure 1.1: Users exploit the Web's Programmability through a Reactive Entity in the Web

Chapter 2

Related Work

In this chapter we will give a brief overview over Information Systems in the Web and their associated Web services as the main drivers for interoperability between Information Systems. We then introduce the Event-Condition-Action(ECA) paradigm, an event-driven approach to impose reactivity. And finally we analyze existing rule languages and engines that exploit the ECA paradigm.

2.1 Data and Functionality Providers on the Web

Information Systems on the Web and their services are also called Web Information Systems and their Web services. Web Information Systems often provide access to internal resources through their services, such as documents or objects, which can be identified and addressed and thus fall into the category of Web Resources. The internal set of resources and their relations form the Web Resource's Information Space[27]. The term service in the context of the Web is somewhat ambiguous and there have been a lot of completely different approaches to offer services on the Web, some of the latest used in cloud computing are Platform as a service (PaaS), Software as a service (SaaS) and Infrastructure as a service (IaaS). The term Web Service (capitalized word Service) commonly stands for Web service based on Simple Object Access Protocol (SOAP) communication[5], which has been adopted and developed extensively by both research and the industry and is often referred to "WS-*" Web Services. With the advent of the Representational State Transfer (REST) architectural paradigm, the understanding of the term service in the Web has undergone a slight generalization so that the term Web service (lower-case word service) is not anymore bound to SOAP, but describes services as interfaces for communication between applications over a network[37]. We will point out main research areas on service-orientation within the Web and show how they make the Web programmable.

Execution of programs on remote computers has always been a strong research area, ever since computer started to exist. After the coinage of the term World Wide Web[6], the Web has become a synonym for Berners-Lee's vision of a global information universe. The adoption of remote program execution through the Web followed immediately; computers sitting in the Web waiting for a request in order to execute some application logic and return an answer. Similar to this concept is the encapsulation of functionality into services[32] in order to offer them to other applications, which is called Service-Oriented Architecture (SOA)[33]. Applying SOA to an existing Information System means splitting it into smaller loosely-coupled pieces (services), which then have to communicate with each other over proper interfaces. Well described server-side service interfaces, meant for application to application communication are called Web APIs.

But the term Web API not only comprises server-side interfaces but also client-sided ones (e.g. the browser), since they are also interfaces to programmability of the Web. Proper interfaces do not only provide robustness, it also allows the reuse of functionality. Moreover these services can be offered to other applications and also to the Web, thus allowing others to access certain functionality or large parts of the Information Systems behind the services. All nodes in the Web are stand-alone entities, which offer services of some sort, be it a webpage, pure data, real-time measurements or functionality of some sort. This makes the Web itself a Service-Oriented Architecture and all the services within it are naturally Web services. It is because of its advantages, that SOA has received a great deal of attention and has been widely adopted throughout the Web. This led to an increasing number of Web accessible services and their compositions, the so called Mashups. An empirical study[24] on a directory, which the researchers of the paper call the "[...] *most active Web APIs and mashups collection*", and statistical data taken from this directory (depicted in Figure 2.1) seem to underline a growing popularity, at least in terms of published services within this particular directory.




Figure 2.1: Number of registered APIs in the ProgrammableWeb directory by date

2.1.1 Accessing services on the Web

An early adoption of the service concept to computers were the Remote Procedure Calls (RPC)[8]. Through RPC a piece of code can be executed on a different machine, other than the one which is calling the procedure. It is basically an inter-process communication and doesn't necessarily require the Web nor distributed computers. RPC also found its use in grid computing[40] and through this, opened doors into the field of distributed computation. The RPC paradigm isn't bound to certain technologies and thus, has been implemented in a lot of different programming languages. These implementations were tightly bound to the respective language that was used, which resulted in incompatibility among them. It became necessary to enhance RPC's in order to get cross platform compatibility. The abstraction of RPC through the Extensible Markup Language (XML)[11], called XML-RPC, made it easier to achieve compatibility between services that base on different technologies.

Since XML-RPC was held relatively simple but received a lot of attention, it was further enhanced. Together with additional proposed functionality, XML-RPC was the base for Simple Object Access Protocol (SOAP)[10]. SOAP is accompanied by the Web Service Description Language (WSDL)[15] which is used to describe the interfaces to SOAP services, the Web Services. With SOAP and WSDL a client of the service can issue a request for the WSDL information of the service and retrieves all interface specifications he requires in order to issue a call to the actual service. The service specifications are then usually incorporated into the existing Information System as if it is a local function call. SOAP has found its applicability

in business applications[5] and was enhanced with a lot of industrial standards, also called the "WS-*" specifications, e.g. WS-Addressing, WS-Policy, WS-Security and many more.

Another initiative that aimed for eased communication between different platform is the one for the Common Object Request Broker Architecture (CORBA)[22]. As the name already suggest it is an object-oriented approach and it allows the access to whole objects over a network. CORBA relies on its communication layer, the Object Request Broker (ORB), which forms the basis of its architecture. The platform-specific ORBs provide the communication abstraction, which free the application from platform dependencies. Similar to SOAP's WSDL, CORBA has its Interface Definition Language (IDL) to provide information about the accessible objects. An object is instantiated by an application and the interface to this instance is offered through the ORB. Another application attached to the ORB can then access all public variables, data structures and functions of this object. This means not only remote access to variables and data structures, but also remote function invocation as seen in RPCs. For programming languages that are not object-oriented, this behaviour has to be simulated, which can be technically difficult and become a tedious task. CORBA enables communication between applications written in different programming languages, no matter whether they run on the same physical device or another one in the network or even the Web. With the Internet Inter-ORB Protocol (IIOP) it is also possible to connect ORB's over the Web. Through this, the offered objects can become services in the Web, but they are shielded by the ORB and only accessible over it.

2.1.2 Web Resources become Services on the Web

All the afore mentioned approaches require a specific protocol and as a consequence are incompatible with each other. For this reason and its simplicity, an architectural style has gained popularity which frees Information Systems and their services from this constraint: Representational State Transfer (REST)[19]. REST concentrates on the roles of components and on constraints upon interactions between them. An important architectural constraint is that all communication is stateless, which means for a client-server communication, no state is stored on the server. Therefore all informations required for a single interaction need to be provided within one request. This allows for the definition of simple and well-defined interfaces, since responses are not bound to a certain session state. Services within the Web that adhere to the REST architecture are called RESTful Web services. RESTful Web services provide access to data and functionality of grouped Web Resources, which can be identified via Uniform Resource Identifiers (URI)[26]. In the Semantic Web[7], a Web Resource is anything in the Web that can be identified, addressed and handled. Historically this started with documents and went over objects to abstract concepts, such as operators of equations. Simple access to Web services without communication overhead and negotiation before using it, increased REST's popularity and made it spread into more application fields. RESTful Web services are often implemented using the HTTP request methods GET, POST, PUT and DELETE which allows for all the necessary operations to create, update, read and delete Web Resources. By using HTTP, the protocol on which the World Wide Web bases, a wide range of Information Systems should be able to communicate with the service. There is for example the upcoming concept of the Web of Things[23], which aims to incorporate smart things (e.g. tagged things, sensor measurements, device controllers, etc.) into the Web through REST interfaces. Even though the nature of such things is usually compatible with the REST architectural constraints, incompatible standards and protocols were used by different manufacturers. Therefore REST brings advantages into the context of smart things connected to the Web.

2.1.3 Composing Services in the Web

Webpages emerged into dynamic sites on the Web through the upcoming of scripting languages to control the browser and the webpage itself. With all their server-sided infrastructure in the background they became literally applications, with more or less functionality and persistence on a server. These Web Applications (Web Apps) became even more responsive with the upcoming of asynchronous calls from the browser to the server, which allows to load data into the current webpage while the user is interacting with it. Those asynchronous calls are requests to services, which act as the Web API to the Web App, which sits on the server. For server-side Web APIs this means that these services can be accessed from other entities in the Web than just browsers, which eases application to application communication. Often the model behind a Web App can be controlled without the Web App itself, depending on how fat the server-side and how thin the client-side is. Imagine not going to the Google webpage anymore to issue a search and manually crawling through the results, but you have your own application doing it for you and processing the results instantly. There is a trend of Web App providers to publish their Web API in order to allow easy access to it. This has lead to an increasing number of Web App Mashups in the past few years.

Mashups combine Web APIs of more than one service in the Web in a new site. Simple services from different sources can be combined into more powerful ones, which can in turn again be composed, and so on. These service compositions assemble data and services in a novel way which provides a new perspective on the data. Ever since services were accessible in a more or less convenient way, Mashups have been developed as well. One of the first Web service Mashups[35], was invented in the same year after Google Maps came up in 2005. It was a webpage that displayed Craigslist's rental houses on a Google Map. At that time no Web API was available, to provide easy access to those two services. But there was value to be observed from anybody being able to create a Mashup through publicly available services, because this leads to new ideas and an increase of popularity in all incorporated Web services. Such Mashups are often a read-only and fixed wirings of different Web APIs that provide a new perspective on specific data. Some recent Mashup examples are:

- Wifi and Plugs: MapBox, Google Docs and Import.io API's used to display where Wi-Fi and plugs are available in London.
- MapLight: GovTrack.us and OpenSecrets API's used to combine political results with financial contributions, in order to show how capital contributions to certain campaigns influence voting.
- Shared Count: Facebook, LinkedIn, Pinterest and Twitter API's used to display informations about how well spread a URL is on social media sites.

But also a number of studies[14][25][38][41] made efforts towards personalized Mashups, where users are capable of choosing what and how to link in order to enhance Web Resources according to their needs. These flexible Mashup applications often provide methods to access user-specific functionality within external Web Apps, which makes them even more user-centered and customizable.

2.1.4 Subscribing to Web Resources

There is another type of service in the Web which is about the opposite of the afore mentioned approaches in terms of the data flow. It is the concept of push notifications on state changes, which is a recent research area. There are some manifestations of this model for server to browser

communication, such as Comet[16] or Server-Sent Events¹. The concept is called Webhooks and introduces instant delivery of data whenever it gets available, compared to the need of actively requesting a service to deliver it. Webhooks are URIs, which point to a Web Resource, which accepts the data delivered to it. Within the publish/subscribe paradigm[18], such asynchronous delivery of data is referred to as events, since it depicts the appearance of new data. Webhooks are callbacks that can be placed by a Web service provider at a remote Web API, informing a distinct event delivering Web Resource about the interest in the promised events. Both parties are a sort of Web service, since the Webhook providers accept the data delivered to their URI and the Webhook receiver accepts URIs and offers to send the data. PubSubHubbub² is an open server-to-server publish/subscribe protocol that uses Webhooks for servers to announce their interest in updates from other servers. Only through such push notifications a reactive system can be real-time reactive through instant event detection.

2.1.5 Towards Simple Access and Communication

With JavaScript's success as browser scripting language and recently also as server-side programming language, JavaScript Object Notification (JSON), as an alternative to XML, has become popular for data representation and communication throughout the Web. It is also because of its human-readable format and often simple parsing into data structures of existing programming languages. There is a notable trend towards RESTful services in the Web that offer JSON communication. They benefit from simple but fully capable interfaces and easy to debug human-readable communication, which eases integration into other applications, along with low communication volume. Together with client- and server-side Web APIs the Web becomes ever more programmable.

2.2 Reactivity through Event-Condition-Action Rules

In this chapter we have so far shown research in different areas, which lead towards a programmable Web. As a result of this research, it is getting easier to compose and orchestrate Information Systems, but reactivity needs to be programmed specifically by experts and general approaches are only available in specific domains. Several studies[3][12][13][28][30] have been made on reactivity. They point out Event-Condition-Action (ECA) rules as an intuitive way to impose reactivity on Information Systems. As the name already suggests it bases on an Event-Driven Architecture (EDA) consist of three parts:

- Event: An event identifier for the to be detected event
- Condition: Expressions to be evaluated to determine whether the action section should be triggered
- Action: A set of instructions that complete the reactive behaviour

Several different rule languages have been developed for different domains. We will give a brief overview over research related to our vision, reactivity in the Web. During our research, apart from identifying the key properties of different rule languages, we analyzed them with respect to a certain use case, in order to determine their applicability for our research goal. The use case's ECA rule is:

¹<http://dev.w3.org/html5/eventsource/>

²<http://code.google.com/p/pubsubhubbub/>

- Event: Receipt of an Email
- Condition: Assert a certain sender email address
- Action: Store it remotely via a Web API

We also tried to get access to existing rule engines for each rule language, since we aim to build our model as well as our own reference implementation on top of existing work.

2.2.1 Rule Languages & Rule Engines

The Resource Description Framework (RDF) is a collection of specifications to model informations in the Semantic Web. Papamarkos et al. (2004) published an ECA language for RDF: RDF Triggering Language (RDFTL). It was designed to react on insertion and deletion events within RDF repositories and as an action propagate the changes through related resources and impose actions on the local repository. RDFTL bases on RDF resources which need to run engines in order to react on the rules. These engines retrieve events, detect changes and communicate them as events to other engines, while actions are executed on local repositories. Through distributed engines, RDF resources are made reactive. We envision an engine that orchestrates the Web, rather than relying on other Information Systems to incorporate our model. But still their research provides important insights on reactivity through ECA rules.

The rule language XChange[31] emerged from the Reasoning on the Web with Rules and Semantics (REWERSE) project[36], which took place from 2004 to 2008. It was designed to track changes in dynamic Web Resources and add reactive behaviour in a way that such changes influence other dynamic resources. XChange incorporates the vision of distributed, event exchanging rule engines. Those rule engines execute actions on local data or issue new events. The local-only actions oppose our vision to orchestrate heterogeneous Web Resources through reactive behaviour, as does the RDFTL. The use case applicability study was promising but access to a reference implementation of an engine, in order to enhance it with our vision, could not be gained. Still the thorough research done with the language XChange holds valuable concepts, especially in terms of temporal event composition.

JSON Rules[21] was introduced 2008 as a language to react on specific Document Object Model (DOM) tree states of a webpage and as reactive behaviour control the browser and also the DOM again. The incorporation of script function calls into the action part of the language allows the abstraction of eventually complex action behaviour. This feature influenced our concept as it allows for different levels of complexity to be offered and regards the different levels of expertise possible users have. JSON Rules is bound to DOM tree events and actions, while we aim to react on any events happening in Web Resources and also execute actions on them.

The Rule Markup Language (RuleML)[9] is a language written in XML and aims to standardize many different types of rules. Reaction RuleML[30] is an enhancement of the existing standard by reactive rules. Reaction RuleML subsumes:

- Complex Event Processing (CEP)
- Knowledge Representation (KR) calculi
- Event-Condition-Action (ECA) rules
- Production (CA) rules
- Trigger (EA) rules

Reaction RuleML represents thorough research for a language to describe virtually any type of reactivity. Together with the expression in XML it doesn't score with readability, but provides a way to define a multitude of rule types and ensures their interchangeability between different Information Systems. Since our vision does not require interchangeable rules, we chose an internal JSON representation, inspired by JSON Rules, for our rules to have more important properties, such as human-readability, simple parsing and efficient storage. A notable system that relies on RuleML is Rule Responder[29]. Rule Responder connects different types of heterogeneous rule engines together over the Mule open-source Enterprise Service Bus (ESB) which acts as a communication middleware to exchange rules expressed in RuleML. The introduction of a communication layer between Information Systems is the same concept as in CORBA.

A recent research outcome (Windley, 2011) is the Kinetic Rule Language (KRL)[42] together with the Kinetic Rules Engine (KRE). It was invented to impose reactivity to the Web and incorporates many different event origins and action resources. The language is based on a declarative syntax, enriched with imperative elements. An interesting feature is the incorporation of the browser into the architecture, which bridges the gap between the user's browser and the centralized KRE. Either a user can install a browser plugin which will communicate with the KRE, or a webpage provider can include a library in order to get events from accesses to the webpage. Through this, events can be raised from the browser and actions can also be executed in it. The KRL fits very well into our concept and only a few reasons kept us from realizing our reference implementation on top of the KRE, such as:

- complexity required to maintain states with a declarative syntax
- system footprint of the KRE
- Perl, a procedural programming language, as basis of the KRE

The concept of the KRL is promising in terms of orchestrating the Web through reactivity. But we made the decision towards a light weighted reference system, using an event-driven programming language that supports Event-Driven Architecture natively. Another important key property of our envisioned conceptual model, which diverges from the KRE architecture, is the abstraction of state maintenance into action dispatching modules, rather than incorporating them into the language.

2.2.2 Overview over existing Rule Languages and their Engines

Table 2.1 gives an overview of the key properties of existing rule languages and their key properties for our research:

- Event Origin: Resource type from where the events originate.
- Distributed: Whether the language is laid out to run on a centralized or distributed architecture. All examined rule languages that support distributed architectures can as well run on a centralized architecture.
- Action Resource: Resource type on which actions are executed.
- Accessible Engine: Lists accessible engine reference implementations for the language.
- Applicability to our concept: Names the main difference to our envisioned concept.

Rule languages that support a distributed architecture require engines to be deployed on sites in order for them to be reactive, since actions are only imposed locally. This is not service-oriented in terms of external services and does not attempt to orchestrate the Web's heterogeneous services

in the action part of rules. It seems common for ECA rules to only invoke actions on local systems, even though the KRL goes into the direction of accessing remote systems too.

Other examined Rule Engines are the Object-Oriented Java Deductive Reasoning Engine for the Web (OO jDrew), Prova, and Drools Fusion. They are all implemented in Java and have their own rules syntax which is more or less closely related to Java, with inline Java code, except for Prova which uses a Prolog like syntax. Some of them have a heavy system footprint because they base on communication middlewares or required frameworks to be run in, such as Drools Fusion which requires the Java graphical user interface eclipse. Our decision not to use any of these existing approaches for a reference implementation was, because all of the examined research only provides limited support for our concept or has a heavy system footprint, caused by the communication layer such as the JBoss ESB. We envision a scalable system, which is event-driven from the application layer, and allows the orchestration of heterogeneous Web Resources. These resources are already available and accessible and we do not need to alter them in order to impose reactivity to the Web. Such a system does not require a messaging middleware because the Web itself acts as the communication channel to receive events and dispatch actions.

2.2.3 Complex Event Processing

An important research area in the field of Event-Driven Architectures is Complex Event Processing (CEP)[4] and deals with event composition, also called Complex Event Detection (CED)[2][34]. It is the research for methods to detect predefined event relations and also temporal patterns, over different streams of data or events. This topic has received a lot of attention for active databases[1][20][43] and was picked up again in the context of Event-Driven Architectures. With CED atomic and composite events are successively aggregated into higher-order events, regarding temporal constraints. The event-driven architecture of reference systems allows the processing of large amounts of data and assemble compositions in real-time. In our conceptual model we envision a CED engine that detects complex event patterns and assembles them into higher-order events. These complex events can then be detected by the ECA rules engine to dispatch appropriate actions.

Language	Event Origin	Distributed	Action Resource	Accessible Engine	Applicability to our concept
RDFTL	RDF Repository Changes	Yes	(Local) RDF Repository	-	Only Web sites with engines are reactive
XChange	Web Resources	Yes	Local Resources	-	Actions in remote Web Resources missing
JSON Rules	DOM Events	No	Browser / DOM	-	Only Browser / DOM Events
RuleML	Web Resources	Yes	Local Resource	(OO) jDrew, Prova, Rule Responder	Complex Syntax
KRL	Web Resources	No	Local & Remote Web Resources	KRE	User-specific Web App functionality missing

Table 2.1: Key Properties of existing Rule Languages

Chapter 3

Conceptual Model for Reactive Information Systems and their Services

The challenges and opportunities arising with the growth of the Web in terms of volume and complexity inspired our research towards the reactive Web. Therefore our starting point were the studies of related work in the context of reactivity on the Web, event composition and programmability of the Web. In the last chapter, we pointed out how they received a lot of attention and provide powerful tools to orchestrate the rapidly growing Web. By combining existing research in these fields we developed a conceptual model, which allows to impose smart reactivity to any Information Space, not only the Web. Even though our initial set of Information Spaces was thought to consist of Web Resources, our model is applicable to any Information System whose Information Space can be accessed and altered over interfaces, i.e. services. Thus we introduce our conceptual model for reactive Information Systems and their services in this chapter.

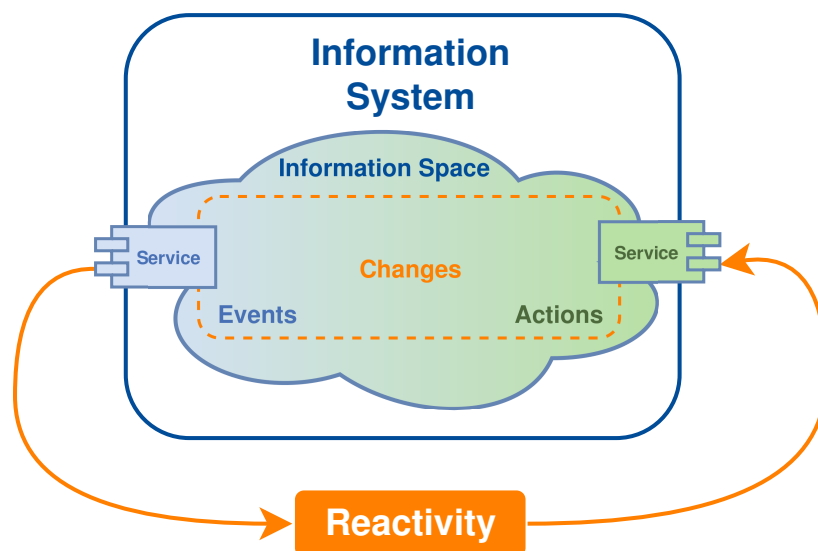


Figure 3.1: Reactivity imposed on Information Systems and their Information Spaces over Services

Data changes within an Information System can be detected and imposed from the outside,

if appropriate interfaces to the services exist. We model the detection of data changes as events, and the imposition of such changes as actions, as shown in Figure 3.1. Through this we are able to introduce an event-based model that is capable to detect events and react on behalf of them by imposing actions on any Information Space, be it the event origin or another one. A more precise distinction of the required modules for such a reactivity imposing entity is displayed in Figure 3.2, and each module is introduced in this chapter.

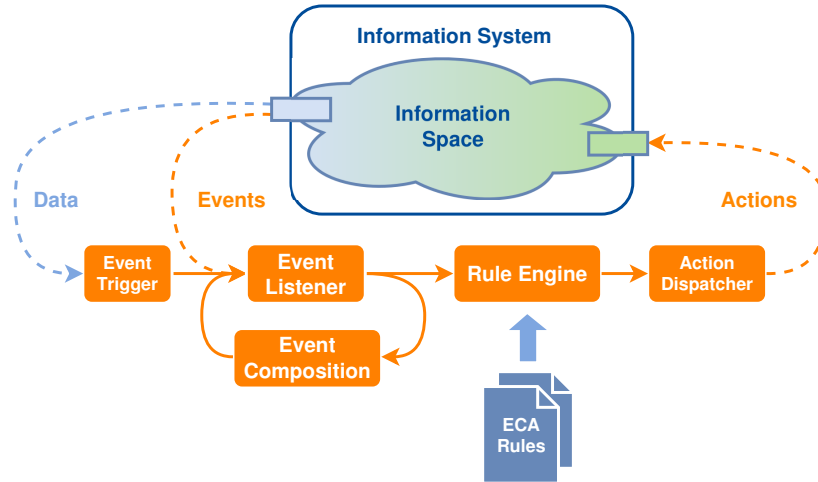


Figure 3.2: Conceptual Model for Reactive Information Systems and their Services

3.1 From Physical Events to Virtual Events

We introduced the Web of Things in the last chapter, where smart devices gain access to the Web. It is based on the Internet of Things which dealt with the incorporation of sensor networks into the Internet on the network level. Such sensors bring the physical world directly into the virtual world. This transformation pictures an important difference between physical and virtual events. In physics, and in particular relativity, an event indicates a physical situation or occurrence, located at a specific point in space and time. While physical events correspond to a physical situation which is located at a specific point in space and time, virtual events primarily consist of implicit parameters, i.e. a name and occurrence time. These virtual events can be anywhere within Information Systems at any point in time and thus their actual location differs most certainly from its occurrence location. But virtual events also have explicit parameters which correspond to all available information about the event, such as the origin. As soon as events are transformed into the virtual world, the afore mentioned location information is transformed into explicit event parameters. Therefore wherever a virtual event is, it has a name, an occurrence time and most likely some explicit parameters attached to it. If the virtual event is of a physical nature, it has a physical location, if it is of a virtual nature, it is likely associated with a virtual origin. Since in our model events are changes in data of an Information Space, they can be virtually anything, e.g. physical measurements, changes on a static webpage, changes of the object behind a Web service or also a login attempt.

3.2 Capturing Events from Information Systems

The optimal case for an event-driven system which requires events from a remote Information System is, that events are triggered within the remote Information Systems and then immediately communicated to interested parties, such as our envisioned reactivity imposing system. But our research has shown that such Information Systems are often passive and rarely provide ways for external systems to announce interest in changes of their data. For example in the context of Web Resources, many of them provide access to their data over services, but do not actively communicate changes to interested parties. This is where the upcoming concept of Webhooks comes into play. Because of effectivity, it is essential for Information Systems to push event notifications to external systems, instead of letting them poll for events. Only through such pushed events it is possible to have real-time reactivity without high costs of continuously polling for changes over all Information Spaces. We also need to take passive Information Systems into account which do not push events to external systems, therefore we need to incorporate polling for changes into our model. Wherever an Information System is not capable to provide events to external Information Systems, we can still read all the accessible data and detect changes in it, model them as events and feed them into our model. In our model the polling for changes is incorporated in the Event Trigger modules. Those are flexible modules that have the proper tools to access any Information System service and therefore its Information Space and are capable of identifying changes in the data. For example the World Wide Web, as envisioned by Tim Berners-Lee[6], is an information universe of interlinked documents, that a user can browse through. Through our model, we can pull changes in the data on the World Wide Web, i.e. document changes, and turn them into events. These events which are derived from changes in the data of Information System are then fed into the Event Listener. The Event Listener also pulls events directly from the Information Systems that offers service functions which represent events but still need to be requested actively, e.g. new mail in inbox.

3.3 Event Pattern Detection

Traditional ECA systems only react on single events, but this might often not be enough to detect meaningful situations. Primitive events occur at a point in time (e.g. a press down mouse button event). When they are composed (e.g. the latter event with a release mouse button event), they turn into a composite event which is more complex and also has a duration. Such a temporal event composition yields the chance to detect meaningful situations out of primitive events and react on them. This is why there is a trend towards the detection of complex event patterns, as we have pointed out in the last chapter. CEP could be incorporated into the rules of the Rule Engine, which then reacts on event patterns, but this opposes our vision of a successively growing complexity of composite events that are defined on top of each other and fed back into the Event Listener. Thus in our model an Event Composition module composes events into more complex events according to CED definitions. It is a very active research field, it has seen interesting studies[2][34] and outcomes¹ that could be incorporated into our model. Such an event composing service systems works loosely coupled and could be realized by any suitable system, as described in [39].

¹such as <http://drools.jboss.org/drools-fusion.html>

3.4 Imposing Reactivity to Information Spaces

In the last chapter we gave an introduction into reactivity and the ECA paradigm as an approach to achieve it. So far we have introduced the first part of our model which provides the foundation of an Event-Driven Architecture. What we also need is a module that translates events into actions on Information Systems. Almost all existing ECA system's actions write on the local Information Space which opposes our vision of the orchestration of different Information Systems in order to impose reactivity on top or between them. For that reason we introduce the Action Dispatcher modules which are located right behind the Rule Engine in terms of the data flow and complete the reactivity flow between heterogeneous Information Systems. Action Dispatcher modules are an important part of our model because they allow flexible coupling with Information System services, much like the Event Trigger modules do. Event Trigger and Action Dispatcher modules are communication abstractions to services of Information Systems, that allow us to deal with their heterogeneity in terms of communication. The Information Space of an Information System is not limited to internal data, but can also refer to a coupling with other devices and the sensing and controlling of it. Thinking of the Web of Things this can also include an Action Dispatcher that has access to an Information System which controls devices and thus is capable of turning down the heating in a house, as shown in Figure 3.3.




Figure 3.3: Information Systems providing access to the Web of Things over their Services

So far we defined all the modules to access Information Systems over their services, therefore we are able to define a Rule Engine that orchestrates them in a reactive way. We have shown in the last chapter that ECA rules consist of three parts; an event to be recognized, conditions to be evaluated on the event and actions to be executed if an event triggers the rule through valid conditions.

Chapter 4

Use Cases for Reactive Information Systems

We have so far introduced a conceptual model for reactive information systems and their services. Through our model we are able to react on events happening in Information Systems and dispatch changes to it or any other accessible Information Systems. In this chapter we will give examples of what use cases can be realized through our model.

4.1 Reacting on changes in the World Wide Web

Many documents of the World Wide Web are dynamic in the way that they change over time. Some might change in an interval of a few minutes (e.g. news), while others change much slower, such as knowledge sites. To detect such changes an Information System in the Web needs to keep track of the document history and trigger events if there is a change. In our model this would be realized by an Event Trigger which monitors a certain set of Information Spaces and triggers events as soon as differences are detected. A user will then set up a rule that checks if the changes are related to a certain category of interest, such as sports or a certain observed site and how big those changes are. If for example a Wikipedia article changed in more than 20% of its content this will be a reason for a moderator to look at it and should be entered as a task in the next free slot of his calendar.

4.2 Enhance existing Web Applications

Web applications, such as webmails, social networks or Content Management System (CMS) are widely spread and used by a large number of Internet users. But often users or developers miss some features or interoperability with other web applications, which would result in enhanced functionality and also in less work. Features of that kind could also include data and functionality from other Web Resource on the web. This would require Web Applications to communicate together and to grab data or impose functionality upon each other. A lot of enhancements will not be implemented by the Web Applications themselves because they are very specific to a small number of their users. With a reactive Information Systems, users and developers could realize such features on their own.

4.2.1 Enrich CMS Posts with Additional Information

Every new post to a Content Management System can be modeled as an event. In the case that a user would like to enrich such a CMS with knowledge from a remote resource, reactivity in the Web can be used to do it. Enhancing an existing CMS can be realized by a rule that evaluates new posts and checks whether knowledge tags are included in the post. Whenever there are tags included within the post, the reactive entity will enrich the post with additional knowledge to these tags from a remote Web Resource of the user's choice.




Figure 4.1: Enrich CMS Post with Remote Knowledge Data

4.2.2 Workflow Automation

Within such Web Applications, users often have very specific workflows. And because workflows always start with an event, they are predestined to be automated by a reactive entity. As an example for workflow automation, course and student exercise submission administration can be taken care of by a reactive entity. Whenever a new semester starts, the reactive entity will detect this through one of its rules and command an action dispatcher to set up infrastructures for courses. This can also include grabbing course data from an official webpage and including it into the infrastructure, thus eliminating the need for manual data copy tasks.

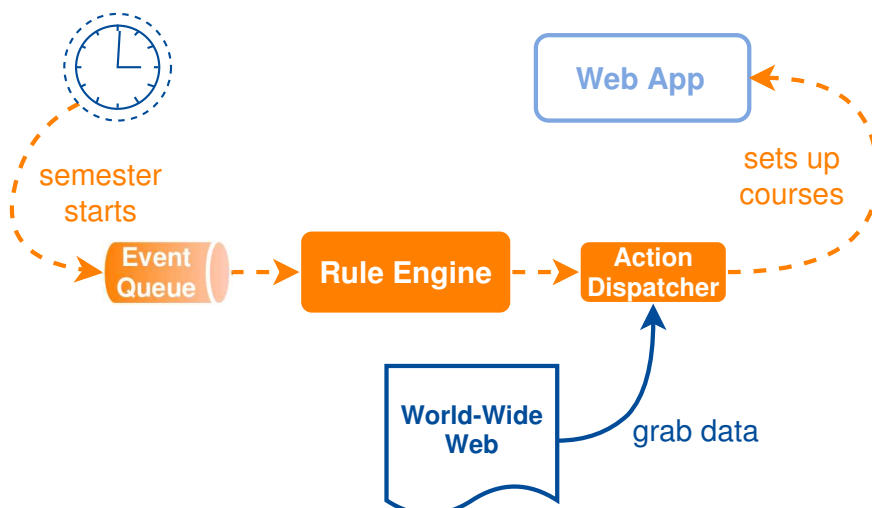


Figure 4.2: Create course resources at semester start

After setting up the semester courses, the reactive system is ready to process new student registrations for these courses. It automatically associates students into the afore mentioned

infrastructures and sets up additional infrastructure such as an exercise submission container. Whenever a course tutor submits a new exercise to the course resource, the system will detect this and spread this information to the students, together with a deadline. The students are expected to submit their exercise solutions before the deadline, to the exercise submission container, which was created reactively for them.

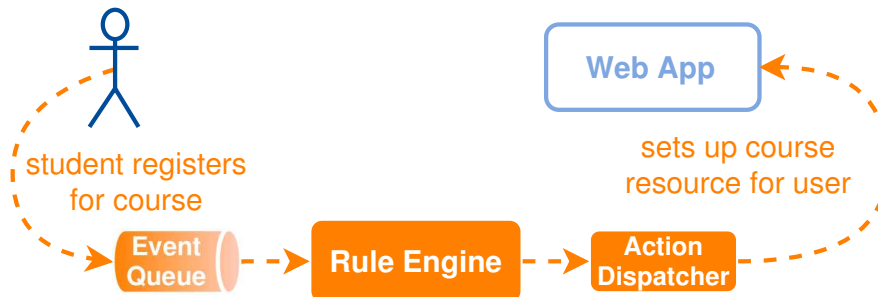


Figure 4.3: Create course resource for registered student

A certain amount of time before the deadline, e.g. one day, the reactive entity will detect the deadline and process events that depict the current exercise submission status per student. If the system detects students who haven't uploaded their exercises yet, it will notify them about the deadline. This is an additional service that gives students the chance to react on a missed exercise submission deadline. As soon as the deadline passed, the system will revoke write-rights to the exercise submission container and therefore disallow submissions which are too late.




Figure 4.4: Notify student before exercise due date

4.3 Service Functionality and Availability Checking

Services offered through the Web aren't monitored or tested by users or developers from other sites. If they rely on correct functionality or availability they need a way to assert this. It is also possible that an owner of such a service doesn't have the tools to monitor his own services. Whenever such a service isn't working correctly anymore or stops responding, these users or developers need to be able to react on this before it is too late. With a reactive rule in place that evaluates Service Testing results, measurements can be taken early. One action to such a failing service test could be an automatic switching of the utilized service within an action dispatcher, so that from then on it uses one which still works correctly.

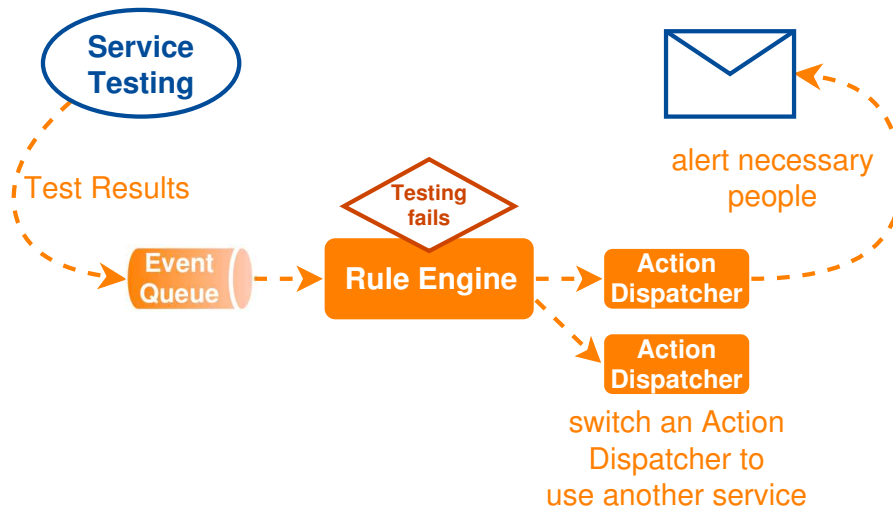


Figure 4.5: Test proper Service Functionality and Availability

4.4 Exploiting the Web of Things

A model for reactive information systems becomes also very interesting in the context of the Web of Things. Through the small connected devices, a lot of sensor data become accessible via the Web and can be used as events to trigger actions. These actions could also be part of the Web of things, if there are Things that offer services. One example of a reactive rule, that has parts in the Web of Things, is that of a server room which has a defective cooling. The increasing temperature eventually causes the servers to shutdown or even fail. Servers in this room could push current state information into a reactive system. The reactive system could take measurements if it detects a certain pattern that will lead to an overheating of all systems. It could inform certain (not so important) servers to gracefully shutdown and additionally inform administrators, who otherwise might miss the shutdown. Even better would be if it would have the power to kick in an additional emergency cooling system to prevent the shutdown of any of the servers.

Another scenario gets more realistic with the increasing number of homes that are connected to the Web. A home or apartment owner has her light controls attached to the Web. The first thing a reactive system could do, is that it detects holidays in the owner's agenda and automatically sets the light control to somewhat reasonable random during her absence. This would make suspicious characters, which are eventually interested in her wealth, think that he's still at home. In combination with another Thing that is connected to the Web and always accompanies people, the cell phone, an even more interesting application scenario can be thought of. The phone would push location information about the owner into the system. Whenever the owner gets close to her house, the reactive system could turn on the light in the entry area and play some music. On a Wednesday evening it could also inform the delivery service that they can deliver the owner's preferred menu when she is home, because the owner is doing this always on a Wednesday evening.

4.5 Multi-Sourced Bad Weather Prediction

There are a lot of different weather services existing in the Web. One has to check several of them in order to get an idea on how likely it is that it will be raining on the journey to work and

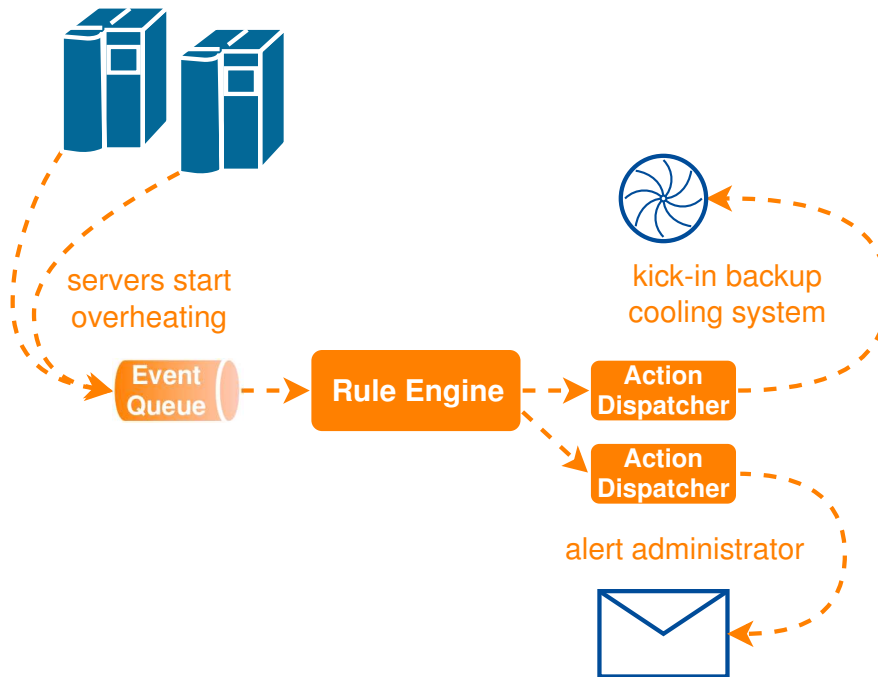


Figure 4.6: Measurements on Server Failure

back. By composing a higher-order event from several weather update events, a user could store a rule which alarms him early in the morning if more than 50% of the weather forecasts expect rain on the way to or back from work.

Chapter 5

Prototype System

We have so far introduced our conceptual model for reactive Information Systems and their Services and some example use cases to point out what would be possible with our model. In this chapter we present our proof of concept prototype system, which has a focus on the Web as its Information Spaces. We will then introduce our ECA rule language, which gives all the necessary power over our prototype system and which can be directly translated into the internal rules representation.

5.1 Architecture

The prototype system is the adoption of our conceptual model for reactive Information Systems and their services to the Web. The Web consists of many Information Systems and because of its Service-Oriented Architecture it can be seen as one large Information System, therefore we can impose reactivity on the Web. Since communication over services in the Web is often latency driven, we came to the conclusion that asynchronous communication and therefore scalability should be attributes our prototype system has to support natively. Another aspect to be regarded for the architectural decision was how the rules are going to be represented internally. We introduced XML and JSON as common ways to communicate data between services on the Web. Both formats represent data in a tree structure, and this is also what we decided to assume for the explicit parameters in the events that will enter our prototype. Together with the requirement of native support for an Event-Driven Architecture (EDA) our decision was to build upon the recent adoption of JavaScript to application development through Node.js¹ and its human-readable JSON communication format.

The prototype system consists of several modules, shown in Figure 5.1, which we are going to introduce within this section:

- **Poller:** Loads Event Trigger modules and forwards events coming from them to the Event Queue. Event Trigger modules poll for changes in the Web and transform them into events.
- **Webhook Listener:** Listens on active Webhook for events and forwards them to the Event Queue.
- **Event Queue:** Buffers events for the case of an overly busy Rule Engine.

¹<http://nodejs.org/>

- **Rules Engine:** Picks an event from the Event Queue whenever there is one and it is idle.
- **User Request Handler:** The user interface modules to administrate Event Triggers, Webhook, Rules and Action Dispatchers.

When started, the prototype system loads persisted Webhook and begins to listen for new events on them. The Rule Engine then loads all persisted rules and for each rule it loads the required Action Dispatchers and notifies the Poller about the new rule, which then in turn loads an Event Trigger if required. The prototype is now up and running and accepts administration requests for Event Triggers, Webhook, Rules and Action Dispatchers. Whenever a rule is created or updated, the Poller and Rule Engine load required Event Triggers or Action Dispatchers.

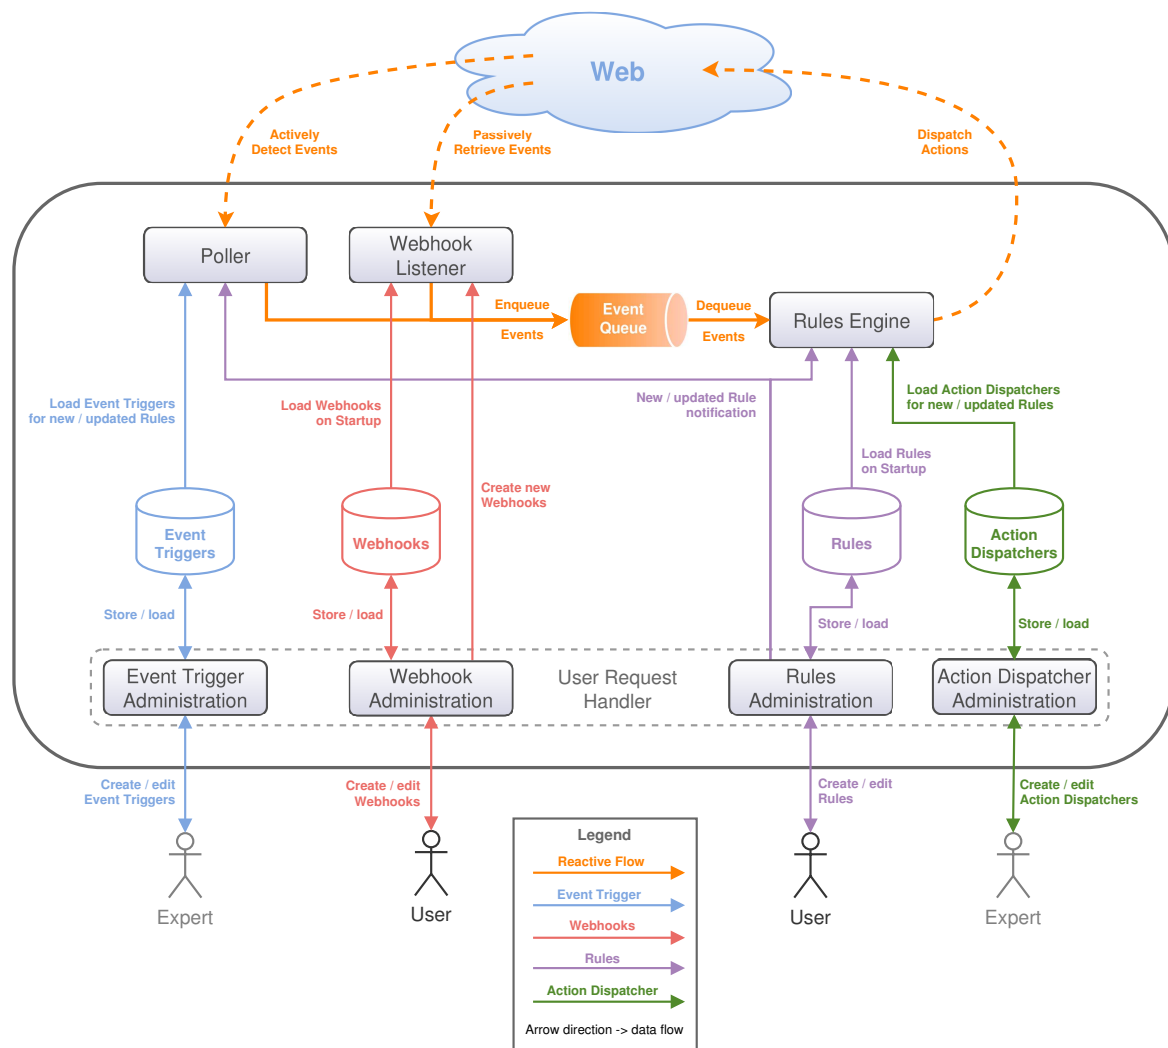


Figure 5.1: Prototype System Architecture

5.1.1 Data Structure for Event Parameters

Events in our prototype system are internally represented as tree structures in JSON format, which builds on only two data structures:

- Objects: Unordered collections of name/value pairs wrapped in curly brackets { }, which can also be implemented as a hash map, dictionary or struct in other languages.
- Arrays: Ordered lists of values wrapped in brackets [], which can also be implemented as a record, vector or list in other languages.

A value can be an object or an array, but also a unicode string, a number, boolean or null. This allows for any arbitrary depth and chaining of the supported data representations. It is a handy feature when we assume a tree structure for events in our prototype system since the selection of nodes in tree structures has been well studied and useful libraries exist. Since all native data structures within JavaScript programming code are already in JSON format, it can easily be marshaled into one string and communicated to other applications without overhead. JSON can be implemented in virtually every programming language, therefore received a lot of attention and is supported by many Web Resources for communication.

5.1.2 Dynamic Code Loading for Event Trigger and Action Dispatcher

During our research we have seen many different Web services with thoroughly different requirements in terms of communication. The cleanest category among them were the Web APIs with their RESTful services. And still, in many cases it is only possible to access data which does not refer to an event. To detect a change in data over time we need to be able to store an earlier request and get the difference to the current request, which could then be transformed into a meaningful event. The derivation of a meaningful event from data can be a complex task which requires certain operations, which underlines the need for powerful Event Trigger modules. The complexity is even bigger for Action Dispatcher modules which alter data and require more complex communication to Web services. For those reasons we made the decision to keep these modules flexible in terms of communication and also to leave it open to expert users to encapsulate complicated logic into them for inexperienced users. Therefore we modeled the Event Trigger and Action Dispatcher to be JavaScript code modules that are created by expert users during runtime and loaded whenever a rule is activated, which needs them. These modules run in a sandbox and got only access to certain JavaScript libraries, which are provided by the owner of the system. Through this it is possible to communicate with RESTful Web services as well as with SOAP Web Services or any other service that can be addressed through JavaScript libraries. Apart from those libraries there are two other important functions offered in both type of modules:

- log: Will store log entries on a per rule base, wherever the instruction is met during execution of a module. The log can be seen by the user who chose the module to be part of a rule.
- pushEvent: This is the important part for Event Trigger modules, since events are only through this function pushed into the prototype system. For Action Dispatcher modules this provides the possibility for loopback events.

Only functions which are attached to the exports property of the module are later accessible from the outside and can be selected as Event Trigger or Action Dispatcher. The function arguments of, from outside visible, functions are identified by the according User Request Handler module and the user will be requested to provide values for all of them in order to activate the Event Trigger or Action Dispatcher. For Action Dispatchers it is also possible to use event property selectors as arguments and thus allows the passing of event data to the Action Dispatchers. Through the pushEvent function in the global scope of the modules, events can be pushed into the prototype system.

By using the expressiveness of JavaScript and some of its libraries, it is possible to access a large part of the existing Information Systems and transform changes in their Information Spaces into events and also to impose changes onto them as part of actions. The power that can be expressed in those code modules needs to be controlled, and can't be granted to anybody, thus we shield it through user access control, thus only allowing trusted users to write Event Trigger and Action Dispatcher code.

5.1.3 Retrieving Events

In the last chapters, we always put emphasis on the two different ways how events can be retrieved from Information Systems, i.e. actively pulling events, or passively retrieving them. Even though some Information Systems offer the access to services that offer data that corresponds to events and can instantly be forwarded into the system, we have seen that there is need for the derivation of events from changes in data on the Web, therefore we need the Event Trigger modules. But still, our vision is that of an optimal real-time reactive Web which means that all possible events are offered by all Information Systems. An interested remote entity could then just announce interest in a certain kind of events over a Webhook and retrieves them in real-time. For that reason we laid out our architecture for such Webhook, but still offer the Event Trigger modules to poll for events in the semi-static World Wide Web. During prototype testing we focused mainly on server-sided Web APIs, but we also generated events from the browser and pushed them to our prototype system. This was achieved with a library included in a sample webpage that pushed events to a Webhook of our prototype. Since modern browsers support geo locating, we decided to let the client browser push the the current position of the device to the Webhook.

Polling with Event Triggers

As we have pointed out before, Event Trigger modules are JavaScript code modules with access to a set of predefined libraries, which offer the flexibility of JavaScript programming. The Poller loads those Event Trigger modules whenever they are required as part of a rule. The user of the Event Trigger can choose a starting point and an interval for the polling to take place. An example Web service which offers polling for events is the Email Yak² Web API, which responds with new emails when requested. The code required to request the new mails from this service and forward them into the prototype system is quite short and shown in Listing 5.1, but for other examples it can get more complex. For better readability the code is written in CoffeeScript³. Only expert users are expected to store such a piece of code in our prototype, which enables inexperienced users to simply choose the "EmailYak -> newMail" Event Trigger for their rule. A great opportunity to access data from webpages via a Web API is Import.io⁴. By browsing through the Web with their browser it is possible to select certain parts from a webpage and store them in a mask. Data is instantly extracted from the webpage, using the stored mask, when sending a request to their Web API with the given mask id and the URI. This is a great tool for expert users to predefine desired data on webpages and then produce events out of an Event Trigger whenever there is a change in that data.

²<http://www.emailyak.com/>

³<http://coffeescript.org/>

⁴<https://import.io/>

```

1 url = "https://api.emailyak.com/v1/#{params.apikey}/json/get/new/email/"
2 exports.newMail = () ->
3   needle.get url, ( err, resp, body ) ->
4     if not err and resp.statusCode is 200
5       pushEvent mail for mail in body.Emails

```

Listing 5.1: Event Trigger code to poll Email Yak RESTful Web service for new Mails; written in CoffeeScript

Webhook

As powerful as Webhook are in providing real-time notifications from remote Web Resources, as simple they are to use. In our prototype, users can create as many new Webhook as they choose to. They only need to provide an event name which will be associated to the Webhook. A new Webhook is created in the Webhook Listener, which from then on accepts events posted to it, and the Webhook URI is displayed to the user. Any Web Resource that supports the concept of Webhook (e.g. GitHub⁵) has a place to register the Webhook URI. Whenever a remote Web Resource pushes an event to the Webhook, the user-defined event name is assigned as the implicit parameter of a freshly created internal event, while the whole incoming event body is added as explicit parameters to it. Afterwards it is forwarded to the Event Queue.

5.1.4 Dispatching Actions

Action Dispatchers are JavaScript code modules that can be created during runtime and loaded by the engine whenever a new rule requires them, much as the Event Trigger modules are loaded by the Poller. They need to be created before they can be used in a rule. In our prototype system, Action Dispatchers use a library for HTTP communication which allows them to address a wide range of Web Resources. Action Dispatchers can also push events back into the Event Queue which can be used to chain certain rules together. Since Webhooks are an important part of our vision we also implemented an Action Dispatcher that delivers events to external Webhook. Action Dispatchers need to have functions attached to their `exports` property so that they are visible from the outside and can be selected as actions, such as the `newContent` function in the example Listing 5.2.

5.1.5 ECA Rules in the Rule Engine

While a car engine converts potential energy into mechanical work, our Rule Engine converts events into changes in Information Systems. We have introduced ECA rules as sufficient approach to impose reactivity on systems and adopted the ECA paradigm for our conceptual model. For our prototype this means that the Rule Engine requires user-defined ECA rules which are compared against incoming events. The Rules Administration within the User Request Handler notifies the Rules Engine about new or updated rules from the user, which then in turn loads required Action Dispatcher modules. For each event in the Event Queue, the engine checks it against its stored ECA rules and dispatches actions whenever the event conforms to the rule's condition part. The three parts of an ECA rule have the following requirements in our prototype system:

⁵<https://developer.github.com/Webhook/>

```

1 urlService = 'https://probiner.com/service/'
2
3 requestService = ( args ) ->
4   url = urlService + args.service + '/' + args.method
5   needle.post url, args.data
6
7 exports.newContent = ( companyId, contextId, content ) ->
8   requestService
9     service: 'content'
10    method: 'save'
11    data:
12      companyId: companyId
13      context: contextId
14      text: content

```

Listing 5.2: Action Dispatcher code to store a new content on the ProBinder RESTful Web service; written in CoffeeScript

- **Event name:** Any arbitrary Unicode string, can refer to the name of an Event Trigger or a Webhook, but also to a custom loopback event.
- **Conditions:** Zero or more instructions to be evaluated against an event. Requires a selector for a node in the tree structure of the event, a comparison operator ($<$, $<=$, $>$, $>=$, $==$, $!=$ or *instr*) and a value.
- **Action Dispatchers:** A list of Action Dispatchers to be invoked if all conditions of the given event evaluate to true. We assume that invocations can be expressed using common function invocation syntax (i.e. `actionFunction(param1, param2[, ...])`) in order to dispatch an action.

A valid rule in the internal JSON representation is shown in Listing 5.3, where we used the predefined EmailYak Action Dispatcher to send a mail to an interested person whenever news about soccer are detected.

Parameter Selectors for Events

Tree node selectors for event parameters are used in conditions to select a parameter which is evaluated. The selectors can also be used to pass event parameters as arguments to the Action Dispatchers. Event tree node selectors for Action Dispatcher arguments are defined by wrapping them into curly brackets and prepended with a hash: `"#{ [selector] }"`. Since an existing JavaScript library⁶ is used to find event parameters with selectors, the following selectors are available⁷:

- ***** : Any node
- **T** : A node of type T, where T is one string, number, object, array, boolean, or null
- **T.key** : A node of type T which is the child of an object and is the value its parents key property
- **T:root** : A node of type T which is the root of the JSON document

⁶<https://github.com/harthur/js-select>

⁷Explanations taken from <http://jsonselect.org/>, which is used by js-select

- **T:nth-child(n)** : A node of type T which is the nth child of an array parent
- **T:nth-last-child(n)** : A node of type T which is the nth child of an array parent counting from the end
- **T:first-child** : A node of type T which is the first child of an array parent (equivalent to T:nth-child(1))
- **T:last-child** : A node of type T which is the last child of an array parent (equivalent to T:nth-last-child(1))
- **T:only-child** : A node of type T which is the only child of an array parent
- **T U** : A node of type U with an ancestor of type T
- **T > U** : A node of type U with a parent of type T
- **T ~ U** : A node of type U with a sibling of type T
- **S1, S2** : Any node which matches either selector S1 or S2
- **T:has(S)** : A node of type T which has a child node satisfying the selector S
- **T:val(V)** : A node of type T with a value that is equal to V
- **T:contains(S)** : A node of type T with a string value contains the substring

```

1 {
2   "eventname": "news",
3   "conditions": [
4     {
5       "selector": ".categories",
6       "operator": "instr",
7       "compare": "soccer"
8     }
9   ],
10  "actions": [
11    "EmailYak->sendMail(\"fan@soccer.com\", \"News about soccer!\", \"{ .body }\")"
12  ]
13 }

```

Listing 5.3: Rule Example expressed in JSON

5.2 A Rule Language for the Prototype System

So far, we introduced the internal representation of the ECA rule language used in our prototype system. For human readability and more intuitive writing, they can be transformed into a phrase representation, through which Listing 5.3 would be written as shown in Listing 5.4. Listing 5.4 shows an example phrase of our envisioned rule language where the retrieval of a new mail will be checked for soccer news and, if confirmed, the mail body will be forwarded to an interested person.

```

1 ON news
2 IF ".categories" instr "soccer"
3 DO EmailYak->sendMail("fan@soccer.com","News about soccer!","#{ .body }")

```

Listing 5.4: Example Phrase in Prototype Rule Language

```

1 expression ::= "ON " event " IF " conditions " DO " actions
2 event      ::= word* ("->" word+)?
3 conditions ::= condition (" AND " condition)*
4 condition  ::= string operator "'" string "'"
5 operator   ::= (" < " | " <= " | " > " | " >= " | " == " | " != " | " instr ")
6 actions    ::= action (" , " action)*
7 action     ::= word* "(" (argument (" , " argument)*)? ")"
8 argument   ::= "'" selstring "'"
9 selstring  ::= (word|selector|" ")
10 selector  ::= "#{ " string "}"
11 string    ::= (word|special|" ")*
12 special   ::= [() : . * > ~ , ]
13 word      ::= [A-Za-z0-9_-]+

```

Listing 5.5: Extended Backus-Naur Form of Prototype Rule Language Syntax

5.3 Prototype Use Case Implementations

5.3.1 Webpage Diff

5.3.2 Enhance Existing Web App

5.3.3 Real-time Reactive Messenger

5.3.4 Server Room Overheating

During our research we found a troublesome server room that shows how the Web of Things can be exposed through our model. This server room suffered from a defective cooling system which lead to a drastic increase of temperature in certain circumstances. As a consequence certain server automatically shut themselves down as safety measurements. Eventually, these shutdowns weren't detected immediately by the people that administered these servers, therefore unnecessary downtimes were the result. As a very quick fix to inform certain administrators about the shutdown of their server, we started pingng these servers and pushed the results int

5.4 Web Application Development

5.4.1 Callback Functions & Asynchronous Closures

Often, optimization approaches and programming language concepts require special attention to avoid common pitfalls. When closures are used as asynchronous functions, developers need to be very careful not to end up with race conditions.

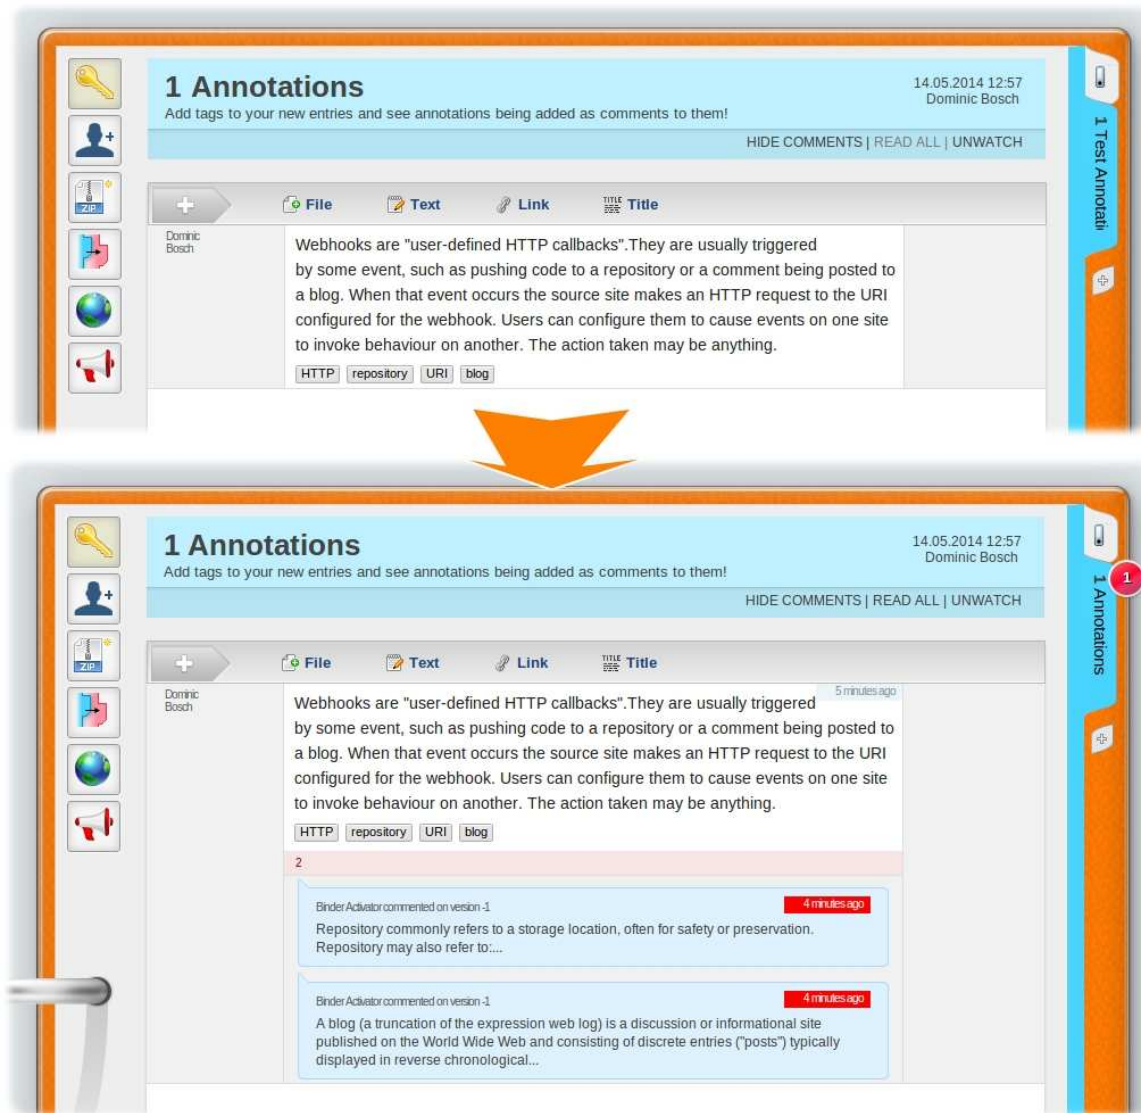


Figure 5.2: UC Binder Annotation

Looking at an example of sequential code execution in Figure 5.3, we see that function execution of fA is halted until function fB is finished. If fB happens to be a latency-driven I/O operation the completion of fA could be deferred for a relatively long time. While the application waits for the completion of the I/O operation, some remaining operations in fA could eventually already be executed without causing any race conditions.




Figure 5.3: Synchronous Function Call

Asynchronous code execution, as shown in Figure 5.4, allows non-blocking and thus scalable applications. Non-blocking operations are a remedy for optimized resource allocation and open up ways to overcome previously described unnecessary resource bindings. Processing any kind of latency-driven I/O operation asynchronously (e.g. filesystem access and socket communication) exploits resources that would otherwise be bound while waiting for completion. Such operations are processed and completed whenever required resources are available.




Figure 5.4: Asynchronous Function Call

Often other operations depend on the completion of asynchronous operations, hence their execution needs to be deferred. This necessary code execution deferral is achieved through the use of callback functions, denoted fCB in Figure 5.4. Any code placed in a callback function, which is assigned to an asynchronous operation, is only executed after the respective asynchronous operation completed. This allows stacking of functions and operations upon each other which automatically results in a flexible and event-driven application.

So far we didn't regard the context for such asynchronous functions. If a function has access to the enclosing context where it was invoked in, it is called a closure. Closures play an important role in ECMAScript[17], which is the base for widely-spread script languages like JavaScript, JScript and ActionScript. Closures in ECMAScript are defined such as they have access to the context of the function they were created in. This is shown in Figure 5.5 where c from fA 's context is accessible from within fB , assuming that fB was created in fA and not only invoked from there. Closures make it necessary for the context of the outer function to survive past its execution so no references are broken. This is labeled "extended context lifetime" in Figure 5.5. Using asynchronous closures it becomes evident, that the context in the invoking function can change while the closure is still computing and eventually referencing the outer context, thus causing race conditions. This will be most obvious in a loop that immediately invokes fB several times, as shown in Figure 5.6. In such a setup c will have different values in the same part of different invocations of fB .

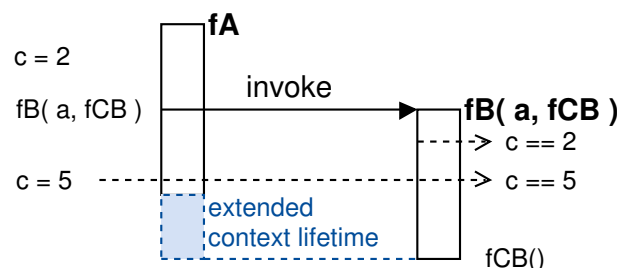


Figure 5.5: Closure Scope and referenced context

Those event-driven context overwrites can be taken care of by shielding the closure from context changes, as shown in Figure 5.7. To shield the closure from context changes, closure




Figure 5.6: Closure context changes in a loop

fB needs to create another closure fC and return it to fA. The argument passed to fB is the context (c in Figure 5.7) that might change but requires to be persistent for one invocation. fC has now c as a fixed context, which can't be overwritten anymore. Now the only thing left is fC needs to be invoked and it will retain the original context. This implementation is necessary when the closure acts as a callback function for asynchronous operations, to preserve the original context in case it is required within the callback function.




Figure 5.7: Closure Context Shielding

An example of how closure contexts can be shielded is shown in the Listing 5.6.

```

1 var fB = function( c ) { // Declare a function...
2   var fC = function( a ) { // ( <-- function to return )
3     console.log( c );
4   };
5   return fC;
6 };
7 for( var c = 0; c < 100; c++ ) {
8   // ... before you assign it to an event happening in the future:
9   var fC = fB( c );
10  setTimeout( fC, 3000 ); // will be executed after the loop ended
11 }

```

Listing 5.6: JavaScript Closure Context Shielding

Chapter 6

Conclusions & Future Work

The practical use case examples for our lightweighted prototype system showed the diversity of examples that can be run in it. As soon as there are services that allow the access to an Information Systems, we are able to orchestrate them, which is a promising insight. But even though the Web gets ever more programmable it is still a tedious task for many Web services to get into communication. Moreover it is a time consuming task to find the proper functionality in Web APIs, or certain functionality needs to be chained in order to get either to meaningful events or actions, or last but not least some functionality is just not accessible. RESTful Web services are a good example for lacking meaningful functionality, since they do not provide complex functionality but more or less read, write, update and delete logic for Web Resources. For those reasons, we believe it was a good decision to encapsulate sometimes complex logic into the Event Trigger and Action Dispatcher modules. Through this, also users that do not have a background in computer science are more likely able to forge their customzied reactive Web.

Many of the notional use case examples in chapter 4 have been implemented in the prototype system, but have brought important insight where it did not work instantly. We found that it is difficult to detect changes on any arbitrary webpage in a useful way. Our first attempt was to use a diff comparison utility, of which the result was hard to process because of the HTML tags. Another approach was to create an object tree from the HTML document and calculate the difference between the last recorded object, which was more promising to detect changes all over a webpage. The simplest approach was to use Import.io to define a part of a webpage to be observed and then detect changes only on this predefined part of the webpage over the Import.io Web API. We believe that there is a lot of value in detecting changes on static Web Resources in a proper way. Future research could mold webpage change detection into some sort of an Event Trigger or even into a novel rule language that uses Web queries in the event part. We were able to realize the enhancement of an existing Web Application in large parts since proper Web APIs provide powerful tools to weild them.

An interesting insight was, that sometimes users wish for activity at a certain point in time instead of event-driven reactivity, which means that the reaction on time events seems to be desirable for users. This could be handled in two ways through our model, either an Event Trigger is created which pushes an event at the exact point in time into the system, or an external Information System is pushing continuously primitive time events into our model over a Webhook. The first option seems to be a bit of an overkill but means lesser event load for the model. The latter option means a certain load depending on how small the event intervals are, but it reflects more the reality where we are used to time events in an interval of one second. Through continous time events it is also easily possible for all users to setup rules for their desired point in time without having to select and parameterize an Event Trigger beforehand.

Our experiences with the prototype system show us that the conceptual model is suitable to impose real-time reactivity on existing Information Systems. We were successful in capturing events from Information Systems, be it by actively pulling them over a service or passively getting it delivered over a Webhook. A somewhat surprising finding was how few Web APIs support Webhooks for real-time notifications to external Information Systems. We envision a future where the whole Web is event-driven and events are directed to any Information System, which is interested in them. This would be the optimal case for effective real-time notifications and thus reactivity on the Web. If the Web APIs remain passive without support for Webhooks, this will cause unnecessary computation and communication cost because of the polling that needs to be done.

A field which turned out to be beyond the scope of this thesis, but was intended to be part of the prototype as well, was the field of CEP. Temporal composition of events should be taken care of in future research in this field since it allows to identify situations out of primitive events. And it also allows to define semantically ever more complex situations out of existing ones. With a growing number of events and their compositions, that are flowing through such a system, it would be useful to have an event relation describing framework on top of them, much like RDF for Web Resources. Also with a growing popularity of a reactivity imposing entity in the Web, other systems will start to deliver events into that entity, which requires detection of new events and information of users about them.

We believe that the Web of Things is a very promising field for such a reactivity imposing entity, but we were not able to study it in depth due to the lacking accessibility of such things. By reducing the interaction in term of event detection and action imposition to RESTful, such as it is used for the Web of Things, it might be possible to incorporate the create, read, update and delete interactions with Web services into the rule language. This would allow for an abstraction away from the Event Trigger and Action Dispatcher modules. It would add stability and a generalization in a way that rules would become more complex to be implemented, but no more expert users are required to implement the code modules. By using the read operation of the RESTful HTTP for webpages, it could be possible to define the detection of changes on them. Since it can be complex to derive an event and express an action with just one call to a service, it might be necessary to allow Web query chaining. Such Web queries could then also be used in the condition section of a rule.

We have seen that the ECA paradigm is very well suitable to impose reactivity to Information Systems and even in an intuitive way for the user to create rules. The more challenging thing was to identify meaningful events that can be added to a rule by the user. We have seen that it is common for existing ECA approaches to impose actions on local data rather than on remote Information Systems over their services. We defined a conceptual model that does not need any Information Systems to understand events but orchestrates it flexibly over its existing services.

Bibliography

- [1] Raman Adaikkalavan and Sharma Chakravarthy. Event Specification and Processing for Advanced Applications: Generalization and Formalization. In Roland Wagner, Norman Revell, and Günther Pernul, editors, *Database and Expert Systems Applications*, volume 4653 of *Lecture Notes in Computer Science*, pages 369–379. Springer Berlin Heidelberg, 2007.
- [2] Mert Akdere, Uğur Çetintemel, and Nesime Tatbul, Plan-based complex event detection across distributed sources. *Proceedings of the VLDB Endowment*, 1(1):66–77, 2008.
- [3] JoséJúlio Alferes and Ricardo Amador. r 3– A Foundational Ontology for Reactive Rules. In Robert Meersman and Zahir Tari, editors, *On the Move to Meaningful Internet Systems 2007: CoopIS, DOA, ODBASE, GADA, and IS*, volume 4803 of *Lecture Notes in Computer Science*, pages 933–952. Springer Berlin Heidelberg, 2007.
- [4] Darko Anicic, Paul Fodor, Sebastian Rudolph, Roland Stühmer, Nenad Stojanovic, and Rudi Studer. A Rule-Based Language for Complex Event Processing and Reasoning. In Thomas Lukasiewicz Pascal Hitzler, editor, *Web Reasoning and Rule Systems - Fourth International Conference*, volume 6333 of *LNCS*, pages 42–57. Springer, September 2010.
- [5] Alistair P. Barros and Marlon Dumas, The Rise of Web Service Ecosystems. *IT Professional*, 8(5):31–37, 2006.
- [6] Tim Berners-Lee, Robert Cailliau, Jean-François Groff, and Bernd Pollermann, World-Wide Web: The Information Universe. *Electronic Networking: Research, Applications and Policy*, 1(2):74–82, 1992.
- [7] Tim Berners-Lee, James Hendler, Ora Lassila, et al., The semantic web. *Scientific american*, 284(5):28–37, 2001.
- [8] Andrew D. Birrell and Bruce Jay Nelson, Implementing Remote Procedure Calls. *ACM Trans. Comput. Syst.*, 2(1):39–59, February 1984.
- [9] Harold Boley. The RuleML family of web rule languages. In *Principles and Practice of Semantic Web Reasoning*, pages 1–17. Springer, 2006.
- [10] Don Box, David Ehnebuske, Gopal Kakivaya, Andrew Layman, Noah Mendelsohn, Henrik Frystyk Nielsen, Satish Thatte, and Dave Winer. Simple object access protocol (SOAP) 1.1, 2000.
- [11] Tim Bray, Jean Paoli, C Michael Sperberg-McQueen, Eve Maler, and François Yergeau, Extensible markup language (XML). *World Wide Web Consortium Recommendation REC-xml-19980210*. <http://www.w3.org/TR/1998/REC-xml-19980210>, 1998.
- [12] Francois Bry and Paula Iavinia Patranjan, Reactivity on the Web: Paradigms and Applications of the Language XChange. *J. of Web Engineering*, 5:2006, 2005.

-
- [13] François Bry and Michael Eckert. Twelve Theses on Reactive Rules for the Web. In Torsten Grust, Hagen Höpfner, Arantza Illarramendi, Stefan Jablonski, Marco Mesiti, Sascha Müller, Paula-Lavinia Patranjan, Kai-Uwe Sattler, Myra Spiliopoulou, and Jef Wijsen, editors, *Current Trends in Database Technology – EDBT 2006*, volume 4254 of *Lecture Notes in Computer Science*, pages 842–854. Springer Berlin Heidelberg, 2006.
 - [14] Cinzia Cappiello, Maristella Matera, Matteo Picozzi, Gabriele Sprega, Donato Barbagallo, and Chiara Francalanci. DashMash: A Mashup Environment for End User Development. In Sören Auer, Oscar Díaz, and GeorgeA. Papadopoulos, editors, *Web Engineering*, volume 6757 of *Lecture Notes in Computer Science*, pages 152–166. Springer Berlin Heidelberg, 2011.
 - [15] Erik Christensen, Francisco Curbera, Greg Meredith, Sanjiva Weerawarana, et al. Web services description language (WSDL) 1.1, 2001.
 - [16] S. Duquennoy, G. Grimaud, and J.-J. Vandewalle. Consistency and scalability in event notification for embedded Web applications. In *Web Systems Evolution (WSE), 2009 11th IEEE International Symposium on*, pages 89–98, Sept 2009.
 - [17] ECMA International. *Standard ECMA-262 - ECMAScript Language Specification*. 5.1 edition, June 2011.
 - [18] Patrick Th. Eugster, Pascal A. Felber, Rachid Guerraoui, and Anne-Marie Kermarrec, The Many Faces of Publish/Subscribe. *ACM Comput. Surv.*, 35(2):114–131, June 2003.
 - [19] Roy Thomas Fielding. *Architectural styles and the design of network-based software architectures*. PhD thesis, University of California, 2000.
 - [20] N. H. Gehani and H. V. Jagadish. Composite event specification in active databases: Model and implementation. pages 327–338, 1992.
 - [21] Adrian Giurca and Emilian Pascalau, Json rules. *Proceedings of the Proceedings of 4th Knowledge Engineering and Software Engineering, KESE*, 425:7–18, 2008.
 - [22] Object Management Group. *The Common Object Request Broker (CORBA): Architecture and Specification*. Object Management Group, 1995.
 - [23] D. Guinard, V. Trifa, and E. Wilde. A resource oriented architecture for the Web of Things. In *Internet of Things (IOT), 2010*, pages 1–8, Nov 2010.
 - [24] Keman Huang, Yushun Fan, and Wei Tan. An Empirical Study of Programmable Web: A Network Analysis on a Service-Mashup System. In Carole A. Goble, Peter P. Chen, and Jia Zhang, editors, *ICWS*, pages 552–559. IEEE, 2012.
 - [25] Xuanzhe Liu, Yi Hui, Wei Sun, and Haiqi Liang. Towards Service Composition Based on Mashup. In *Services, 2007 IEEE Congress on*, pages 332–339, July 2007.
 - [26] Larry Masinter, Tim Berners-Lee, and Roy T Fielding, Uniform resource identifier (URI): Generic syntax. 2005.
 - [27] Gregory B Newby. Metric multidimensional information space. In *TREC*. Citeseer, 1996.
 - [28] George Papamarkos, Alexandra Poulouvasilis, and Peter T Wood. RDFTL: An event-condition-action language for RDF. In *Proc. of the 3rd International Workshop on Web Dynamics*, 2004.
 - [29] A. Paschke, H. Boley, B. Craig, and A. Kozlenkov. Rule Responder: RuleML-Based Agents for Distributed Collaboration on the Pragmatic Web, 2007.

- [30] Adrian Paschke, Harold Boley, Zhili Zhao, Kia Teymourian, and Tara Athan. Reaction RuleML 1.0: Standardized Semantic Reaction Rules. In Antonis Bikakis and Adrian Giurca, editors, *Rules on the Web: Research and Applications*, volume 7438 of *Lecture Notes in Computer Science*, pages 100–119. Springer Berlin Heidelberg, 2012.
- [31] Paula-lavinia Patranjan. *The Language XChange*. PhD thesis, Ludwig-Maximilians-Universität München, 2005.
- [32] C. Peltz, Web services orchestration and choreography. *Computer*, 36(10):46–52, Oct 2003.
- [33] Randall Perrey and Mark Lycett. Service-oriented architecture. In *Applications and the Internet Workshops, 2003. Proceedings. 2003 Symposium on*, pages 116–119. IEEE, 2003.
- [34] P.R. Pietzuch, B. Shand, and J. Bacon, Composite event detection as a generic middleware extension. *Network, IEEE*, 18(1):44–55, Jan 2004.
- [35] Paul Rademacher. HousingMaps. <http://www.housingmaps.com/>, 2005. Accessed: 2014-6-6.
- [36] REWERSE - Reasoning on the Web with Rules and Semantics. <http://rewerse.net>. Accessed: 2014-05-08.
- [37] Leonard Richardson and Sam Ruby. *RESTful web services*. O'Reilly Media, Inc., 2008.
- [38] Sven Rizzotti and Helmar Burkhart. useKit: A Step Towards the Executable Web 3.0. In *Proceedings of the 19th International Conference on World Wide Web, WWW '10*, pages 1175–1176, New York, NY, USA, 2010. ACM.
- [39] D Robins. Complex event processing. In *Second International Workshop on Education Technology and Computer Science. Wuhan*, 2010.
- [40] Keith Seymour, Hidemoto Nakada, Satoshi Matsuoka, Jack Dongarra, Craig Lee, and Henri Casanova. Overview of GridRPC: A remote procedure call API for Grid computing. In *Grid Computing—GRID 2002*, pages 274–278. Springer, 2002.
- [41] Kathryn T. Stolee, Sebastian Elbaum, and Anita Sarma, Discovering how end-user programmers and their communities use public repositories: A study on Yahoo! Pipes. *Information and Software Technology*, 55(7):1289 – 1303, 2013.
- [42] P. Windley. *The Live Web: Building Event-Based Connections in the Cloud*. Cengage Learning PTR, 2011.
- [43] Detlef Zimmer and Rainer Unland. On The Semantics Of Complex Events In Active Database Management Systems. pages 392–399. IEEE Computer Society Press, 1999.

Appendices

UNIVERSITÄT BASEL

PHILOSOPHISCH-NATURWISSENSCHAFTLICHE FAKULTÄT

Declaration on Scientific Integrity

(including a Declaration on Plagiarism and Fraud)

Bachelor's / Master's Thesis (*Please cross out what does not apply*)

Title of Thesis (*Please print in capital letters*):

First Name, Surname (*Please print in capital letters*):

Matriculation No.: _____

I hereby declare that this submission is my own work and that I have fully acknowledged the assistance received in completing this work and that it contains no material that has not been formally acknowledged.

I have mentioned all source materials used and have cited these in accordance with recognised scientific rules.

In addition to this declaration, I am submitting a separate agreement regarding the publication of or public access to this work.

☐ Yes ☐ No

Place, Date: _____

Signature: _____

Please enclose a completed and signed copy of this declaration in your Bachelor's or Master's thesis .

